



## Problem Statement

A mobile network provider is facing a business problem as lot of customers are transitioning to other service providers. This is causing significant loss to the business. The company likes to understand the factors which impact the loss of customers.

## Dataset used



Telecom\_Data.csv

## Variables

Variables	Description
Churn	1 if customer cancelled service, 0 if not
AccountWeeks	number of weeks customer has had active account
ContractRenewal	1 if customer recently renewed contract, 0 if not
DataPlan	1 if customer has data plan, 0 if not
DataUsage	gigabytes of monthly data usage
CustServCalls	number of calls into customer service
DayMins	average daytime minutes per month
DayCalls	average number of daytime calls
MonthlyCharge	average monthly bill
OverageFee	largest overage fee in last 12 months
RoamMins	average number of roaming minutes

## Requirement

- Explore the relationship between the variables.
- Develop a predictive model to predict which of the customers will churn out of the network.



## Coding in R

```
# SVM classification on Customer churn #
```

```
### Loading package - "readr" ###
```

```
library(readr)
```

```
Cust_DF = read_csv("Telecom_Data.csv")
```

```
> head(Cust_DF)
```

	Churn	AccountWeeks	ContractRenewal	DataPlan	DataUsage	CustServCalls	DayMins	DayCalls	MonthlyCharge	OverageFee	RoamMins
1	0	128	1	1	2.7	1	265.1	110	89	9.87	10.0
2	0	107	1	1	3.7	1	161.6	123	82	9.78	13.7
3	0	137	1	0	0.0	0	243.4	114	52	6.06	12.2
4	0	84	0	0	0.0	2	299.4	71	57	3.10	6.6
5	0	75	0	0	0.0	3	166.7	113	41	7.42	10.1
6	0	118	0	0	0.0	0	223.4	98	57	11.03	6.3

```
> colnames(Cust_DF)
```

	"Churn"	"AccountWeeks"	"ContractRenewal"	"DataPlan"	"DataUsage"	"CustServCalls"
[1]	"Churn"	"AccountWeeks"	"ContractRenewal"	"DataPlan"	"DataUsage"	"CustServCalls"
[7]	"DayMins"	"DayCalls"	"MonthlyCharge"	"OverageFee"	"RoamMins"	

```
### checking the structure of the variables ###
```

```
> str(Cust_DF)
```

```
'data.frame': 3333 obs. of 11 variables:
```

	\$ Churn	\$ AccountWeeks	\$ ContractRenewal	\$ DataPlan	\$ DataUsage	\$ CustServCalls	\$ DayMins	\$ DayCalls	\$ MonthlyCharge	\$ OverageFee	\$ RoamMins
	: int	: int	: int	: int	: num	: int	: num	: int	: num	: num	: num
		128 107 137 84 75 118 121 147 117 141 ...	1 1 1 0 0 0 1 0 1 0 ...	1 1 0 0 0 0 1 0 0 1 ...	2.7 3.7 0 0 0 0 2.03 0 0.19 3.02 ...	1 1 0 2 3 0 3 0 1 0 ...	265 162 243 299 167 ...	110 123 114 71 113 98 88 79 97 84 ...	89 82 52 57 41 57 87.3 36 63.9 93.2 ...	9.87 9.78 6.06 3.1 7.42 ...	10 13.7 12.2 6.6 10.1 6.3 7.5 7.1 8.7 11.2 ...

```
### checking spread of the data ##
```

```
> summary(Cust_DF)
```

	Churn	AccountWeeks	ContractRenewal	DataPlan	DataUsage	CustServCalls	DayMins
Min.	:0.0000	Min. : 1.0	Min. :0.0000	Min. :0.0000	Min. :0.0000	Min. :0.000	Min. : 0.0
1st Qu.	:0.0000	1st Qu.: 74.0	1st Qu.:1.0000	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:1.000	1st Qu.:143.7
Median	:0.0000	Median :101.0	Median :1.0000	Median :0.0000	Median :0.0000	Median :1.000	Median :179.4
Mean	:0.1449	Mean :101.1	Mean :0.9031	Mean :0.2766	Mean :0.8165	Mean :1.563	Mean :179.8
3rd Qu.	:0.0000	3rd Qu.:127.0	3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:1.7800	3rd Qu.:2.000	3rd Qu.:216.4
Max.	:1.0000	Max. :243.0	Max. :1.0000	Max. :1.0000	Max. :5.4000	Max. :9.000	Max. :350.8

	DayCalls	MonthlyCharge	OverageFee	RoamMins
Min.	: 0.0	Min. : 14.00	Min. : 0.00	Min. : 0.00
1st Qu.	: 87.0	1st Qu.: 45.00	1st Qu.: 8.33	1st Qu.: 8.50
Median	:101.0	Median : 53.50	Median :10.07	Median :10.30
Mean	:100.4	Mean : 56.31	Mean :10.05	Mean :10.24
3rd Qu.	:114.0	3rd Qu.: 66.20	3rd Qu.:11.77	3rd Qu.:12.10
Max.	:165.0	Max. :111.30	Max. :18.19	Max. :20.00

```
### Loading package - "psych" ###
```

```
library(psych)
```



```
> describe (Cust_DF)
      vars  n  mean  sd median trimmed  mad min  max range skew kurtosis  se
Churn      1 3333  0.14  0.35  0.00    0.06  0.00  0  1.00  1.00  2.02    2.07 0.01
AccountWeeks 2 3333 101.06 39.82 101.00 100.77 40.03 1 243.00 242.00 0.10   -0.11 0.69
ContractRenewal 3 3333  0.90  0.30  1.00    1.00  0.00  0  1.00  1.00 -2.72    5.42 0.01
DataPlan     4 3333  0.28  0.45  0.00    0.22  0.00  0  1.00  1.00  1.00   -1.00 0.01
DataUsage    5 3333  0.82  1.27  0.00    0.58  0.00  0  5.40  5.40  1.27    0.04 0.02
CustServCalls 6 3333  1.56  1.32  1.00    1.42  1.48  0  9.00  9.00  1.09    1.72 0.02
DayMins      7 3333 179.78 54.47 179.40 179.85 53.82 0 350.80 350.80 -0.03   -0.02 0.94
DayCalls    8 3333 100.44 20.07 101.00 100.57 19.27 0 165.00 165.00 -0.11    0.24 0.35
MonthlyCharge 9 3333  56.31 16.43  53.50  55.22 15.57 14 111.30  97.30  0.59   -0.02 0.28
OverageFee   10 3333 10.05  2.54 10.07 10.05  2.55  0  18.19 18.19 -0.02    0.02 0.04
RoamMins     11 3333 10.24  2.79 10.30 10.28  2.67  0  20.00 20.00 -0.24    0.60 0.05
```

```
### Checking missing values ####
```

```
> sum (is.na (Cust_DF))
[1] 0
```

```
#### Observation - No missing values in the dataset. ####
```

```
### Checking duplicated values ###
```

```
> sum (duplicated (Cust_DF))
[1] 0
```

```
### Observation - No duplicated values. ###
```

```
### Checking the count of Customer categories ###
```

```
> table (Cust_DF $Churn)

 0    1
2850 483
```

```
### Loading package - "plyr" ###
```

```
library (plyr)
```

```
> count (Cust_DF $Churn)
  x freq
1 0 2850
2 1  483
```

```
#### Observation - The above results show that the count of customers churned
#### out is 483. ####
```

```
### Loading package - "ggcorrplot" ###
```

```
> library(ggcorrplot)
Loading required package: ggplot2

Attaching package: 'ggplot2'

The following objects are masked from 'package:psych':

    %+%, alpha

Warning messages:
1: package 'ggcorrplot' was built under R version 4.0.5
2: package 'ggplot2' was built under R version 4.0.4
```

```
> corr = round (cor (Cust_DF), 1)
> head (corr)
```

	Churn	AccountWeeks	ContractRenewal	DataPlan	DataUsage	CustServCalls	DayMins	DayCalls	MonthlyCharge
Churn	1.0	0	-0.3	-0.1	-0.1	0.2	0.2	0	0.1
AccountWeeks	0.0	1	0.0	0.0	0.0	0.0	0.0	0	0.0
ContractRenewal	-0.3	0	1.0	0.0	0.0	0.0	0.0	0	0.0
DataPlan	-0.1	0	0.0	1.0	0.9	0.0	0.0	0	0.7
DataUsage	-0.1	0	0.0	0.9	1.0	0.0	0.0	0	0.8
CustServCalls	0.2	0	0.0	0.0	0.0	1.0	0.0	0	0.0


  

	OverageFee	RoamMins
Churn	0.1	0.1
AccountWeeks	0.0	0.0
ContractRenewal	0.0	0.0
DataPlan	0.0	0.0
DataUsage	0.0	0.2
CustServCalls	0.0	0.0

```
### Loading package - "GGally" ###
```

```
> library (GGally)
Registered S3 method overwritten by 'GGally':
  method from
+.gg      ggplot2
```

```
qqcorr (Cust_DF, label = TRUE, label_size = 2.9, hjust = 1, layout.exp = 2)
```





```
#### observation - High correlations are observed between DataPlan & DataUsage,
#### DataPlan & MonthlyCharge, DataUsage & MonthlyCharge. There's moderate
#### correlation between DayMins & MonthlyCharge. ###
```

```
### Changing the data type of the response variable ###
```

```
> class (Cust_DF $Churn)
[1] "integer"
```

```
> Cust_DF $Churn = as.factor (Cust_DF $Churn)  #### Cust_DF $Churn = sapply (Cust_DF $Churn, factor) ####
>
> class (Cust_DF $Churn)
[1] "factor"
>
> str (Cust_DF)
'data.frame':  3333 obs. of  11 variables:
 $ Churn      : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
 $ AccountWeeks : int  128 107 137 84 75 118 121 147 117 141 ...
 $ ContractRenewal: int  1 1 1 0 0 0 1 0 1 0 ...
 $ DataPlan     : int  1 1 0 0 0 0 1 0 0 1 ...
 $ DataUsage    : num  2.7 3.7 0 0 0 0 2.03 0 0.19 3.02 ...
 $ CustServCalls : int  1 1 0 2 3 0 3 0 1 0 ...
 $ DayMins      : num  265 162 243 299 167 ...
 $ DayCalls     : int  110 123 114 71 113 98 88 79 97 84 ...
 $ MonthlyCharge : num  89 82 52 57 41 57 87.3 36 63.9 93.2 ...
 $ OverageFee   : num  9.87 9.78 6.06 3.1 7.42 ...
 $ RoamMins     : num  10 13.7 12.2 6.6 10.1 6.3 7.5 7.1 8.7 11.2 ...
```

```
## Splitting observational data into 70:30 ratio for training/testing model ##
```

```
> sampl_sz = floor (0.7 * nrow (Cust_DF))
> sampl_sz
[1] 2333

> set.seed (123)
>
> train_ind = sample (seq_len (nrow (Cust_DF)), size = sampl_sz)
>
> train_data = Cust_DF [train_ind,]
>
> test_data = Cust_DF [- train_ind,]
>
> dim (train_data)
[1] 2333  11
>
> dim (test_data)
[1] 1000  11
```

```
## Developing SVM Model ##
```

```
### Loading/Installing package - "e1071" ###
```

```
library (e1071)
```



```
> svm.model = svm (Churn ~ ., data = train_data)
>
> svm.model
```

```
Call:
svm(formula = Churn ~ ., data = train_data)
```

```
Parameters:
  SVM-Type:  C-classification
  SVM-Kernel: radial
    cost: 1
```

```
Number of Support Vectors: 663
```

```
> summary (svm.model)
```

```
Call:
svm(formula = Churn ~ ., data = train_data)
```

```
Parameters:
  SVM-Type:  C-classification
  SVM-Kernel: radial
    cost: 1
```

```
Number of Support Vectors: 663
```

```
( 369 294 )
```

```
Number of Classes: 2
```

```
Levels:
0 1
```

```
## Predicting values ##
```

```
> preds = predict (svm.model, test_data)
>
> tab = table (preds, test_data $Churn)
>
> tab
```

```
preds    0    1
    0 849  75
    1   8  68
```

```
## Computing classification Accuracy ##
## Loading package - "caret" ##
```

```
> library (caret)
Loading required package: lattice
Warning message:
package 'caret' was built under R version 4.0.4
```



```
> confusionMatrix (table (preds, test_data $churn))  
Confusion Matrix and Statistics
```

```
preds    0    1  
  0 849  75  
  1   8  68
```

```
          Accuracy : 0.917  
          95% CI   : (0.8981, 0.9334)  
No Information Rate : 0.857  
P-Value [Acc > NIR] : 4.482e-09
```

```
          Kappa    : 0.5792
```

```
McNemar's Test P-Value : 4.342e-13
```

```
          Sensitivity : 0.9907  
          Specificity : 0.4755  
Pos Pred Value : 0.9188  
Neg Pred Value : 0.8947  
Prevalence : 0.8570  
Detection Rate : 0.8490  
Detection Prevalence : 0.9240  
Balanced Accuracy : 0.7331
```

```
'Positive' class : 0
```

```
### Accuracy observed as 91.7%. ###
```

```
> ## Computing Misclassification Error Rate ##  
>  
> 1 - (sum (diag (tab)) / sum (tab))  
[1] 0.083
```

```
### Observation - Misclassification Error rate computed as 8.3%. ###
```

```
## Developing SVM model with kernel parameter as "linear" ##
```

```
> svm_model = svm (Churn ~ ., data = train_data, kernel = "linear")  
>  
> svm_model
```

```
call:  
svm(formula = Churn ~ ., data = train_data, kernel = "linear")
```

```
Parameters:  
  svm-Type: C-classification  
  svm-Kernel: linear  
    cost: 1
```

```
Number of Support Vectors: 879
```

```
>  
> summary (svm_model)
```

```
call:  
svm(formula = Churn ~ ., data = train_data, kernel = "linear")
```

```
Parameters:  
  svm-Type: C-classification  
  svm-Kernel: linear  
    cost: 1
```

```
Number of Support Vectors: 879
```

```
( 539 340 )
```

```
Number of Classes: 2
```

```
Levels:  
 0 1
```



```
### Observation - No. of support vectors increased. ###
```

```
## Predicting values ##
```

```
> preds1 = predict (svm_model, test_data)
>
> tab1 = table (preds1, test_data $churn)
```

```
## Computing classification accuracy ##
```

```
## Loading package - "caret" ##
```

```
> confusionMatrix (table (preds1, test_data $churn))
Confusion Matrix and Statistics
```

```
preds1    0    1
      0 857 143
      1   0   0
```

```
Accuracy : 0.857
95% CI : (0.8338, 0.8781)
No Information Rate : 0.857
P-Value [Acc > NIR] : 0.5223
```

```
Kappa : 0
```

```
Mcnemar's Test P-Value : <2e-16
```

```
Sensitivity : 1.000
Specificity : 0.000
Pos Pred Value : 0.857
Neg Pred Value : NaN
Prevalence : 0.857
Detection Rate : 0.857
Detection Prevalence : 1.000
Balanced Accuracy : 0.500
```

```
'Positive' Class : 0
```

```
### Observation - Accuracy decreased to 85.7% ###
```

```
## Computing Misclassification Error Rate ##
```

```
> 1 - (sum (diag (tab1)) / sum (tab1))
[1] 0.143
```

```
### Observation - Error rate increased to 14.3%. ###
```





```
## Developing SVM model with kernel parameter as "polynomial" ##
```

```
> svm.model.poly = svm (Churn ~ ., data = train_data, kernel = "polynomial")
>
> svm.model.poly
```

```
Call:
svm(formula = Churn ~ ., data = train_data, kernel = "polynomial")
```

```
Parameters:
  SVM-Type:  C-classification
  SVM-Kernel: polynomial
    cost:    1
   degree:   3
  coef.0:    0
```

```
Number of Support Vectors:  558
```

```
> summary (svm.model.poly)
```

```
Call:
svm(formula = Churn ~ ., data = train_data, kernel = "polynomial")
```

```
Parameters:
  SVM-Type:  C-classification
  SVM-Kernel: polynomial
    cost:    1
   degree:   3
  coef.0:    0
```

```
Number of Support Vectors:  558
```

```
( 296 262 )
```

```
Number of Classes:  2
```

```
Levels:
 0 1
```

```
## Predicting on test data ##
```

```
> preds_poly = predict (svm.model.poly, test_data)
>
> tab_poly = table (preds_poly, test_data $Churn)
```



```
## Computing classification Accuracy ##  
## Loading package - "caret" ##
```

```
> confusionMatrix(table(preds_poly, test_data $churn))  
Confusion Matrix and Statistics
```

```
preds_poly  0   1  
0  846  64  
1   11  79
```

```
Accuracy : 0.925  
95% CI : (0.9069, 0.9406)  
No Information Rate : 0.857  
P-Value [Acc > NIR] : 2.041e-11
```

```
Kappa : 0.6381
```

```
McNemar's Test P-Value : 1.920e-09
```

```
Sensitivity : 0.9872  
Specificity : 0.5524  
Pos Pred Value : 0.9297  
Neg Pred Value : 0.8778  
Prevalence : 0.8570  
Detection Rate : 0.8460  
Detection Prevalence : 0.9100  
Balanced Accuracy : 0.7698
```

```
'Positive' Class : 0
```

```
### Observation - Accuracy increased to 92.5%. ###
```

```
## Computing Misclassification Error Rate ##
```

```
> 1 - (sum(diag(tab_poly)) / sum(tab_poly))  
[1] 0.075
```

```
### Observation - Error rate decreased to 7.5%. ###
```

```
## Developing SVM model with kernel parameter as "sigmoid" ##
```

```
> svm.model.sig = svm(Churn ~ ., data = train_data, kernel = "sigmoid")  
>  
> svm.model.sig
```

```
Call:
```

```
svm(formula = Churn ~ ., data = train_data, kernel = "sigmoid")
```

```
Parameters:
```

```
SVM-Type: C-classification  
SVM-Kernel: sigmoid  
cost: 1  
coef.0: 0
```

```
Number of Support Vectors: 631
```



```
> summary (svm.model.sig)
```

```
Call:
svm(formula = Churn ~ ., data = train_data, kernel = "sigmoid")
```

```
Parameters:
  SVM-Type:  C-classification
  SVM-Kernel: sigmoid
    cost:    1
   coef.0:   0
```

```
Number of Support Vectors:  631

( 317 314 )
```

```
Number of classes:  2
```

```
Levels:
 0 1
```

```
## Predicting values ##
```

```
> preds_sig = predict (svm.model.sig, test_data)
>
> tab_sig = table (preds_sig, test_data $Churn)
```

```
## Computing classification accuracy ##
## Loading package - "caret" ##
```

```
> confusionMatrix (table (preds_sig, test_data $Churn))
Confusion Matrix and Statistics
```

```
preds_sig   0   1
           0 779 125
           1  78  18
```

```
          Accuracy : 0.797
          95% CI   : (0.7707, 0.8215)
 No Information Rate : 0.857
 P-Value [Acc > NIR] : 1.000000
```

```
          Kappa   : 0.0404
```

```
McNemar's Test P-value : 0.001244
```

```
          Sensitivity : 0.9090
          Specificity : 0.1259
       Pos Pred Value : 0.8617
       Neg Pred Value : 0.1875
          Prevalence : 0.8570
       Detection Rate : 0.7790
 Detection Prevalence : 0.9040
       Balanced Accuracy : 0.5174
```

```
 'Positive' Class : 0
```

```
### Observation - Accuracy rate is highly dipped to 79.7%. ###
```



```
## Computing Misclassification Error Rate ##
```

```
> 1 - (sum (diag (tab_sig)) / sum (tab_sig))
[1] 0.203
```

```
### Error rate is highly increased to 20.3%. ###
```

```
### Observation - The SVM model with kernel parameter as "polynomial" has
### the least misclassification error rate & highest accuracy. ###
```

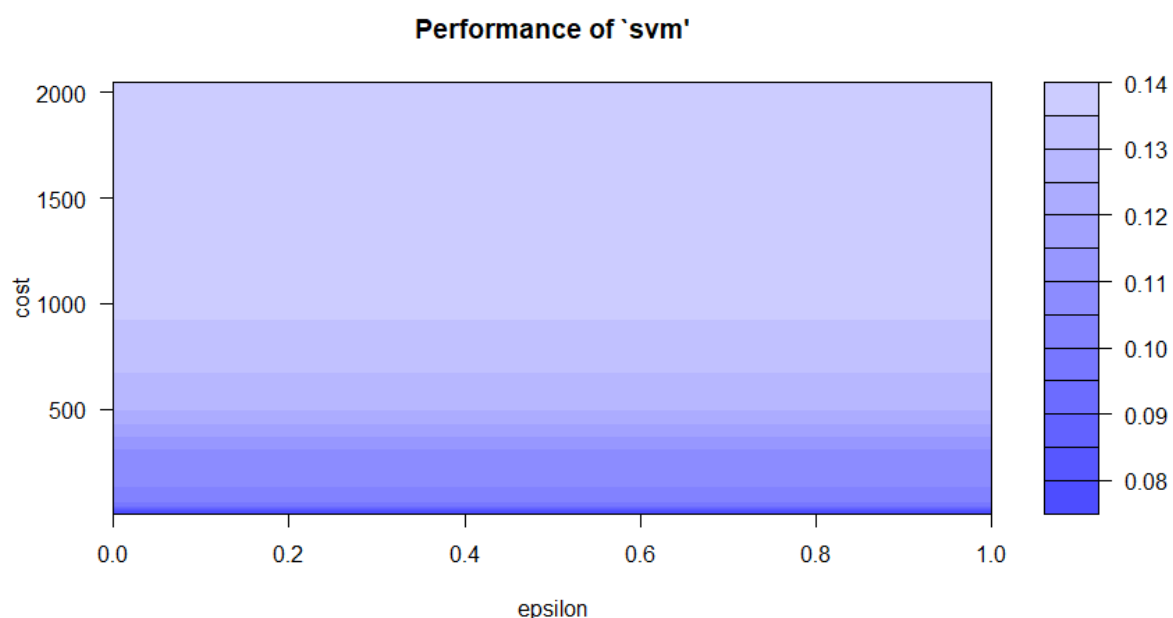
```
## Tuning of model ##
```

```
set.seed (123)
```

```
tmodel = tune (svm, Churn ~ ., data = test_data,
               ranges = list (epsilon = seq (0, 1, 0.1), cost = 2 ^ (2 : 11)))
```

```
## Plotting the model tuned ##
```

```
plot (tmodel)
```



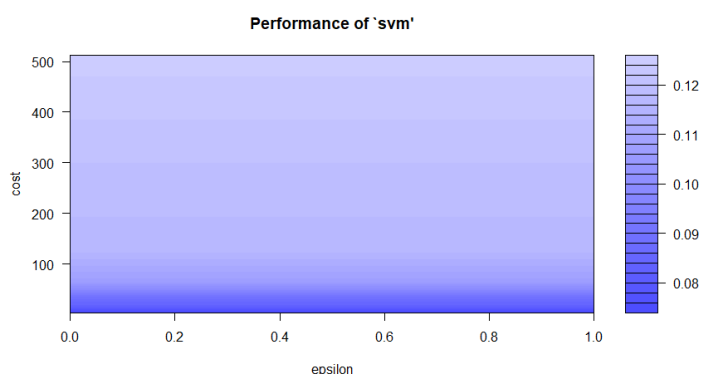
```
### observation - The plot returns performance evaluation of SVM for the
### 2 parameters (cost, epsilon) used. The darker regions towards the bottom
### indicate lower misclassification error. It also suggests that if we
### confine our search till 512 instead of 2048, it will probably make model
### more accurate. ###
```



```
tmodel1 = tune (svm, Churn ~ ., data = test_data,
               ranges = list (epsilon = seq (0, 1, 0.1), cost = 2 ^ (2 : 9)))
```

```
## Plotting the model re-tuned ##
```

```
plot (tmodel1)
```




```
### Observation - The plot returns performance evaluation of SVM for the
### 2 parameters (cost, epsilon) used. The darker regions in the bottom
### indicate lower misclassification error. It also suggests that if we
### confine our search till 200 instead of 512, it will probably make model
### more accurate. ###
```

```
tmodel2 = tune (svm, Churn ~ ., data = test_data,
               ranges = list (epsilon = seq (0, 1, 0.1), cost = 2 ^ (2 : 8)))
```

```
## Plotting the model re-tuned ##
```

```
plot (tmodel2)
```



```
### Observation - The darker regions in the bottom indicate lower
### misclassification error. ###
```



```
> summary (tmodel2)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

```
epsilon cost
0      4
```

- best performance: 0.074

- Detailed performance results:

	epsilon	cost	error	dispersion
1	0.0	4	0.074	0.02547330
2	0.1	4	0.074	0.02547330
3	0.2	4	0.074	0.02547330
4	0.3	4	0.074	0.02547330
5	0.4	4	0.074	0.02547330
6	0.5	4	0.074	0.02547330
7	0.6	4	0.074	0.02547330
8	0.7	4	0.074	0.02547330
9	0.8	4	0.074	0.02547330
10	0.9	4	0.074	0.02547330
11	1.0	4	0.074	0.02547330

```
### Observation - The summary result shows best parameters as epsilon - 0 &
### cost - 4, i.e., 2^2. The best svm classification model can be chosen
### using this summary result. ###
```

```
## Best SVM Model ##
```

```
Best.svm = tmodel2 $best.model
```

```
> summary (Best.svm)
```

Call:

```
best.tune(method = svm, train.x = Churn ~ ., data = test_data, ranges = list(epsilon = seq(0, 1,
0.1), cost = 2^(2:8)))
```

Parameters:

```
SVM-Type: C-classification
SVM-kernel: radial
cost: 4
```

Number of Support Vectors: 279

```
( 171 108 )
```

Number of Classes: 2

Levels:

```
0 1
```

```
### Observation - The summary shows 279 nos. of support vectors with 171, 108,
### nos. of support vectors for each of the 2 classes of the response
### variable. It also shows kernel parameter as "radial"
### & the cost value as 4. ###
```

```
## Predicting on test data ##
```

```
preds_best = predict (Best.svm, test_data)
```

```
tab_Best = table (preds_best, test_data $churn)
```



```
## Computing SVM Classification Accuracy ##  
## Loading package - "caret" ##
```

```
> confusionMatrix(table(preds_best, test_data $Churn))  
Confusion Matrix and Statistics
```

```
preds_best  0   1  
0  854  35  
1    3 108
```

```
      Accuracy : 0.962  
      95% CI   : (0.9482, 0.973)  
No Information Rate : 0.857  
P-Value [Acc > NIR] : < 2.2e-16
```

```
      Kappa : 0.829
```

```
McNemar's Test P-value : 4.934e-07
```

```
      Sensitivity : 0.9965  
      Specificity : 0.7552  
Pos Pred Value : 0.9606  
Neg Pred Value : 0.9730  
Prevalence : 0.8570  
Detection Rate : 0.8540  
Detection Prevalence : 0.8890  
Balanced Accuracy : 0.8759
```

```
'Positive' Class : 0
```

```
### Observation - SVM model Accuracy sharply increased to 96.2%. ###
```

```
## Computing Misclassification Error Rate ##
```

```
> 1 - (sum(diag(tab_Best)) / sum(tab_Best))  
[1] 0.038
```

```
### Misclassification Error rate has remarkably decreased to 3.8%. ###
```

---